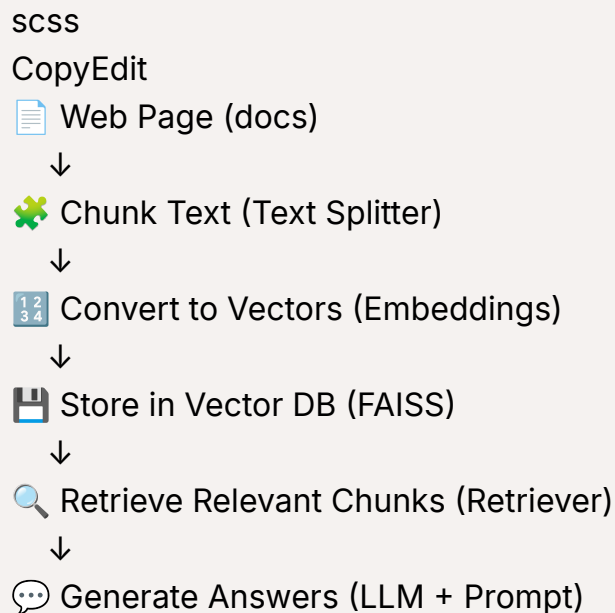


RAG Gen AI App using LangChain + OpenAI

Steps:

- Scraping data from a webpage
- Chunking the data
- Embedding the chunks
- Storing in a vector database (FAISS)
- Creating a retriever
- Passing relevant documents to an LLM to generate answers

High-Level Structure

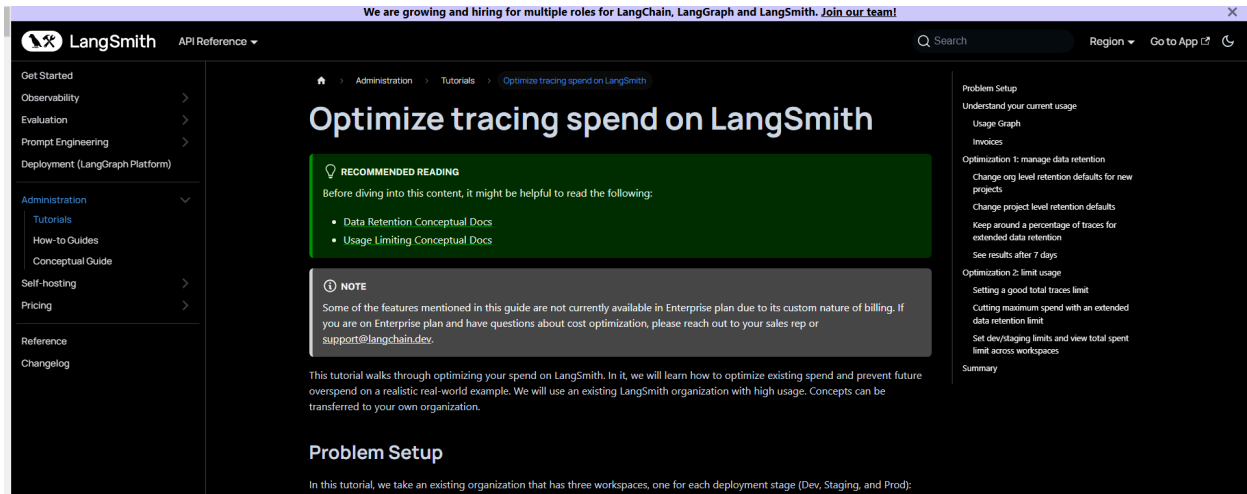


1. Setup API Keys and Environment

✅👉 Scrapes the page content and loads it as documents into `docs`.

Webpage:

https://docs.smith.langchain.com/administration/tutorials/manage_spend



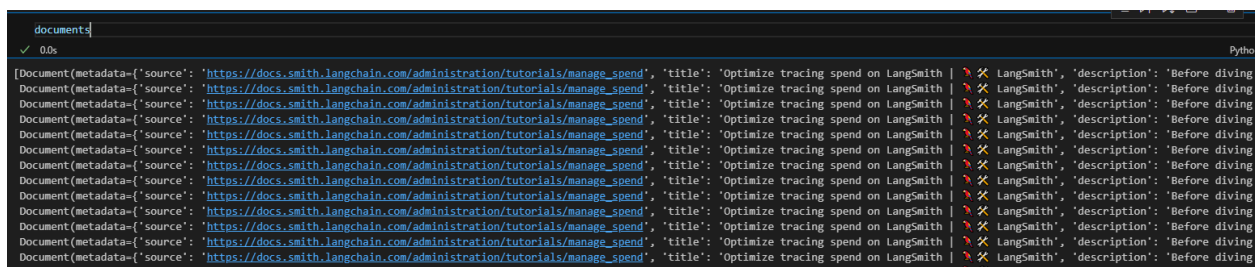
3. ✂ Split Text into Chunks

```
from langchain_text_splitters import RecursiveCharacterTextSplitter
```

```
text_splitter = RecursiveCharacterTextSplitter(chunk_size=1000, chunk_overlap=200)
```

```
documents = text_splitter.split_documents(docs)
```

✅ Breaks long text into manageable parts (chunks), with overlap to preserve context.



4. Generate Embeddings (Vectors)

```
from langchain_openai import OpenAIEmbeddings

embeddings = OpenAIEmbeddings()
```

✓ Converts each chunk of text into a vector (number list) using OpenAI's embedding model.

5. Store Vectors in Vector DB (FAISS)

```
from langchain_community.vectorstores import FAISS

vectorstoredb = FAISS.from_documents(documents, embeddings)
```

✓ Stores all vectors in a searchable FAISS database.

6. Query Similar Texts (Optional Check)

```
query = "LangSmith has two usage limits: total traces and extended"
result = vectorstoredb.similarity_search(query)
print(result[0].page_content)
```

Output:

'use usage limits to prevent future overspend. LangSmith has two usage limits: total traces and extended retention traces. These correspond to the two metrics we've been tracking on our usage graph. We can use these in tandem to have granular control over spend. To set limits, we navigate back to Settings → Usage and Billing → Usage configuration. There is a table at the bottom of the page that lets you set usage limits per workspace. For each workspace, the two limits appear, along with a cost estimate: Let's start by setting limits on our production usage, since that is where the majority of spend comes from. Set

ing a good total traces limit\u200bPicking the right "total traces" limit depends on the expected load of traces that you will send to LangSmith. You should\u200bnclearly think about your assumptions before setting a limit.For example:Current Load: Our gen AI application is called between 1.2-1.5 times per second, and each API request has a trace associated with it,'

✅ Checks how well your query finds the most relevant document.

7. 🧠 Load LLM

```
from langchain_openai import ChatOpenAI

llm = ChatOpenAI(model="gpt-4o")
```

✅ Loads OpenAI's latest GPT model.

8. 💬 Define Prompt and Response Chain

```
## Retrieval Chain, Document chain

from langchain.chains.combine_documents import create_stuff_documents_chain
from langchain_core.prompts import ChatPromptTemplate

prompt=ChatPromptTemplate.from_template(
    """
    Answer the following question based only on the provided context:
    <context>
    {context}
    </context>
    """
)
```

```
"""
```

```
)
```

```
document_chain=create_stuff_documents_chain(llm,prompt)
```

```
document_chain
```

```
RunnableBinding(bound=RunnableBinding(bound=RunnableAssign(mapper={
  context: RunnableLambda(format_docs)
}), config={'run_name': 'format_inputs'}))
| ChatPromptTemplate(input_variables=['context'], messages=[HumanMessagePromptTemplate(prompt=PromptTemplate(input_variables=['context'], output_parser=StrOutputParser(), config={'run_name': 'stuff_documents_chain'})
```

ChatPromptTemplate. from_template

- **What it is:** This method is a simpler way to create a prompt template, especially useful when you have a single, coherent block of text that includes placeholders. While it's a "chat" prompt template, it's more flexible and often used when you're injecting large chunks of information (like retrieved documents) directly into the prompt.
- **Why it's used now:** In your current RAG (Retrieval-Augmented Generation) example, you're going to inject a `<context>` block directly into the prompt. It's often easier to define this entire structure as one string template rather than breaking it down into separate system/user messages, especially when the context itself is dynamic. It implicitly creates a single "user" message with the full template content.
- **Structure:** It takes a single string that contains placeholders (like `{context}` and `{input}`).

In Simple Words:

- **from_messages** is like giving the AI a script with different characters (system, user). You tell it who says what.
- **from_template** is like giving the AI a fill-in-the-blanks letter. It just fills in the missing parts of a single message.

- You choose based on whether you want to strictly define roles in a conversation or just inject dynamic content into a single, larger instruction.

`<context> {context} </context>`

For the LLM: These tags act as clear visual markers. They tell the LLM, "Hey, everything between these `<context>` and `</context>` tags is the specific information you should use to answer the upcoming question. Do *not* use your general knowledge for this part." This helps prevent the LLM from "hallucinating" (making up answers) or using outdated information.

What is `create_stuff_documents_chain(...)` ?

How it works:

1. It receives a list of `Document` objects (e.g., from your `retriever`).
2. It takes the `page_content` from each `Document` and concatenates them into one long string.
3. It then takes this combined string and injects it into the `{context}` placeholder of your `prompt` template.
4. Finally, it sends this fully formed prompt to the `llm` to get an answer.

It's a type of **Document Chain** — a chain that handles:

- Taking retrieved docs
- Formatting them into a prompt
- Getting a smart response from the LLM

Internally:

- When you create a document chain using `create_stuff_documents_chain(...)`
- LangChain **automatically wraps** the output using a `StrOutputParser()`

✨ Why is it called "stuff"?

LangChain offers multiple **strategies** to feed documents to LLMs:

Strategy	Description
<code>stuff</code>	Stuff all docs into the prompt at once
<code>map_reduce</code>	Summarize each doc, then combine summaries
<code>refine</code>	Add one doc at a time, refining answer progressively

So `create_stuff_documents_chain()` is the simplest — just "stuff it all in."

! When Not to Use Stuffing

If your docs are **too big** to all fit in one prompt (e.g., long PDFs), stuffing doesn't work well.

Instead, use `map_reduce` or `refine` strategies — or a **retriever-based RAG chain**.

```
from langchain_core.documents import Document
document_chain.invoke({
    "input": "LangSmith has two usage limits: total traces and extended",
    "context": [Document(page_content="LangSmith has two usage limits: total
traces and extended traces. These correspond to the two metrics we've been
tracking on our usage graph. ")]
})
```

```
'LangSmith has two usage limits: total traces and extended traces. These correspond to the two metrics tracked on their usage graph.'
```

9. Set Up Retriever + Final Retrieval Chain


```
retriever = vectorstoredb.as_retriever()
```

```
from langchain.chains import create_retrieval_chain  
retrieval_chain = create_retrieval_chain(retriever, document_chain)
```

✓ Retriever finds the most relevant docs for a query and passes them to the LLM.

What is a **Retriever** ?

In LangChain, a **Retriever** is like a very smart librarian. Its job is to take a natural language question (your query) and go into the **vectorstoredb** (the library) to find the most relevant "books" or "pages" (document chunks) that might contain the answer. It doesn't answer the question itself; it just finds the relevant context.

```
retrieval_chain = create_retrieval_chain(retriever, document_chain)
```

This is where you combine the "Retrieval" part and the "Generation" (or "Document Answering") part into one seamless process.

- **create_retrieval_chain** : It automates the process of connecting the document finding step with the answer generation step.
- **retriever (Your Librarian)**: This is the first component you pass. It's responsible for finding the relevant documents.
- **document_chain (Your Expert and Report Writer)**: This is the second component you pass. Remember, your **document_chain** already contains:
 1. Your **ChatPromptTemplate** (with the `<context>{context}</context>` format).
 2. Your **ChatOpenAI** LLM.
 3. The automatically added **StrOutputParser**
- **Stage 1: Retrieval**: It first takes the user's **input** question and passes it to the **retriever**. The **retriever** goes to the **vectorstoredb** and fetches the most relevant **Document** objects.
- **Stage 2: Document Answering (Generation)**: It then takes these *retrieved Document* objects and the original **input** question, and passes them both to your

`document_chain`. The `document_chain` then uses the `prompt` to formulate a query with the retrieved context for the LLM, gets the LLM's answer, and extracts the string content.

✓ `retrieval_chain` Internal Flow of this chain:

User Input (query) →



Retriever.fetch_relevant_docs(query) →



Docs injected into prompt as {context} →



LLM answers the query using the documents →



Final response returned

← END 10. Invoke the App

```
response = retrieval_chain.invoke({"input": "LangSmith has two usage limits: total traces and extended"})
print(response['answer'])
```

✓ LLM answers the question using the most relevant documents only.



response (The Team's Result): The `response` you get back from `retrieval_chain.invoke()` is a **Python dictionary**.

It contains key pieces of information from the entire process:

- `'input'`: Your original question.

- `'context'`: The actual `Document` objects that the `retriever` found and passed to the `document_chain`. (This is very useful for debugging! You can see *what* information the LLM actually received.)
- `'answer'`: This is the final, plain string answer generated by your LLM based on the retrieved context and your question.

```
print(response['answer'])
```

(Getting the Final Answer): This line simply accesses the `'answer'` key from the `response` dictionary and prints its value, which is the clear, concise answer from the LLM.

```
response
```

Output:

```
{'input': 'LangSmith has two usage limits: total traces and extended',
 'context': [Document(page_content='use usage limits to prevent future overs
pend.LangSmith has two usage limits: total tra,
...
metadata={'source': 'https://docs.smith.langchain.com/tutorials/Administrato
rs/manage_spend', 'title': 'Optimize tracing spend on LangSmith | 🦜🔧 LangS
mith', 'description': 'Before diving into this content, it might be helpful to read t
he following:', 'language': 'en'}),
```

```
Document(page_content="more than 10% of traces.notelf you want to keep a
subset of traces for longer
```

```
...
...
```

```
'answer': 'To set usage limits in LangSmith, navigate to Settings → Usage a
nd Billing → Usage configuration. There, you will find a table at the bottom o
f the page that allows you to set usage limits per workspace. The two types of
limits you can set are total traces and extended retention traces, each
with an associated cost estimate.'
```

✓ Full Process Visualization:

Your Custom Documents (Vector DB)



Retriever (finds best chunks)



Document Chain (adds prompt)



LLM (generates answer)



Final Output